

Manuel_Pasieka-Wisconsin_Breast_Cancer

February 20, 2019

1 Introducción

El enfoque de esta práctica crear un clasificador para el dataset 'Wisconsin Breast Cancer' Dataset utilizando un 'isolation forest'.

Los datos están disponibles en: <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>

2 Resultados

El mejor clasificador consigue una f1 score de 0.68 para los casos benévolos y 0.81 para los casos malignos.

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import sklearn
from sklearn.ensemble import IsolationForest
```

3 Análisis descriptivo

El dataset contiene 569 ejemplos de pruebas histológicas con 357 ejemplos benévolo y 212 ejemplos malignos.

Hay 30 características numéricas, sin valores missing.

```
In [2]: T = pd.read_csv('data/data.csv')
T.head()
```

```
Out[2]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
0	842302	M	17.99	10.38	122.80	1001.0	0
1	842517	M	20.57	17.77	132.90	1326.0	0
2	84300903	M	19.69	21.25	130.00	1203.0	0
3	84348301	M	11.42	20.38	77.58	386.1	0
4	84358402	M	20.29	14.34	135.10	1297.0	0

	smoothness_worst	compactness_worst	concavity_worst	concave points_worst	symmetry
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	

2	0.1444	0.4245	0.4504	0.2430
3	0.2098	0.8663	0.6869	0.2575
4	0.1374	0.2050	0.4000	0.1625

In [3]: T.describe().T

```
Out [3]:
```

	count	mean	std	min	25%
id	569.0	3.037183e+07	1.250206e+08	8670.000000	869218.000000
radius_mean	569.0	1.412729e+01	3.524049e+00	6.981000	11.700000
texture_mean	569.0	1.928965e+01	4.301036e+00	9.710000	16.170000
perimeter_mean	569.0	9.196903e+01	2.429898e+01	43.790000	75.170000
area_mean	569.0	6.548891e+02	3.519141e+02	143.500000	420.300000
smoothness_mean	569.0	9.636028e-02	1.406413e-02	0.052630	0.086370
compactness_mean	569.0	1.043410e-01	5.281276e-02	0.019380	0.064920
concavity_mean	569.0	8.879932e-02	7.971981e-02	0.000000	0.029560
concave points_mean	569.0	4.891915e-02	3.880284e-02	0.000000	0.020310
symmetry_mean	569.0	1.811619e-01	2.741428e-02	0.106000	0.161900
fractal_dimension_mean	569.0	6.279761e-02	7.060363e-03	0.049960	0.057700
radius_se	569.0	4.051721e-01	2.773127e-01	0.111500	0.232400
texture_se	569.0	1.216853e+00	5.516484e-01	0.360200	0.833900
perimeter_se	569.0	2.866059e+00	2.021855e+00	0.757000	1.606000
area_se	569.0	4.033708e+01	4.549101e+01	6.802000	17.850000
smoothness_se	569.0	7.040979e-03	3.002518e-03	0.001713	0.005169
compactness_se	569.0	2.547814e-02	1.790818e-02	0.002252	0.013080
concavity_se	569.0	3.189372e-02	3.018606e-02	0.000000	0.015090
concave points_se	569.0	1.179614e-02	6.170285e-03	0.000000	0.007638
symmetry_se	569.0	2.054230e-02	8.266372e-03	0.007882	0.015160
fractal_dimension_se	569.0	3.794904e-03	2.646071e-03	0.000895	0.002248
radius_worst	569.0	1.626919e+01	4.833242e+00	7.930000	13.010000
texture_worst	569.0	2.567722e+01	6.146258e+00	12.020000	21.080000
perimeter_worst	569.0	1.072612e+02	3.360254e+01	50.410000	84.110000
area_worst	569.0	8.805831e+02	5.693570e+02	185.200000	515.300000
smoothness_worst	569.0	1.323686e-01	2.283243e-02	0.071170	0.116600
compactness_worst	569.0	2.542650e-01	1.573365e-01	0.027290	0.147200
concavity_worst	569.0	2.721885e-01	2.086243e-01	0.000000	0.114500
concave points_worst	569.0	1.146062e-01	6.573234e-02	0.000000	0.064930
symmetry_worst	569.0	2.900756e-01	6.186747e-02	0.156500	0.250400
fractal_dimension_worst	569.0	8.394582e-02	1.806127e-02	0.055040	0.071460

In [4]: # How balanced are the target clases
T['diagnosis'].value_counts()

```
Out [4]: B    357
M    212
Name: diagnosis, dtype: int64
```

In [5]: # Are there any missing datapoints?
T.isna().sum()

```

Out [5]: id                0
         diagnosis         0
         radius_mean      0
         texture_mean     0
         perimeter_mean   0
         area_mean        0
         smoothness_mean  0
         compactness_mean 0
         concavity_mean   0
         concave points_mean 0
         symmetry_mean    0
         fractal_dimension_mean 0
         radius_se        0
         texture_se       0
         perimeter_se     0
         area_se          0
         smoothness_se    0
         compactness_se   0
         concavity_se     0
         concave points_se 0
         symmetry_se      0
         fractal_dimension_se 0
         radius_worst     0
         texture_worst    0
         perimeter_worst  0
         area_worst       0
         smoothness_worst 0
         compactness_worst 0
         concavity_worst  0
         concave points_worst 0
         symmetry_worst   0
         fractal_dimension_worst 0
         dtype: int64

```

4 Preparación de datos

Cambiar el orden de filas al azar.

Quitar la columna 'id'


Primero calcular el valor standardizado para cada característica.

```

In [6]: R = T.sample(frac=1)
         R.drop('id', axis=1, inplace=True)
         Y = R['diagnosis']
         X = R.drop('diagnosis', axis=1)

In [7]: plt.figure(figsize=(18,13))
         sns.heatmap(X.corr(), annot=True)
         plt.show()

```



4.1 Quitar características con una correlación alta

Quitar características con una correlación alta ha creado resultados peores, así he decidido <> hacerlo.

El código aquí está para documentar el intento

```
highly_correlated = ['id', 'radius_mean', 'perimeter_mean', 'radius_worst',
'texture_worst', 'area_worst', 'perimeter_worst', 'area_se', 'radius_se',
'compactness_worst', 'concavity_worst', 'concave points_worst',
'smoothness_worst', 'fractal_dimension_se', 'concavity_se', 'concave points_se',
'concavity_mean', 'concave points_mean'] X.drop(highly_correlated, axis=1,
inplace=True) plt.figure(figsize=(18,13)) sns.heatmap(X.corr(), annot=True)
plt.show()
```

```
In [8]: nX = (X - X.mean())/X.std()
nX.head()
```

```
Out [8]:      radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean  compactness_mean
40      -0.195029      0.532512      -0.238242      -0.261112      -1.048076      -0.811112
212       3.967796     -0.190570       3.972634       5.240230       1.268455       0.811112
287     -0.351099     -1.434456     -0.414792     -0.394952     -1.906288     -1.811112
```

314	-1.569301	-0.160345	-1.558873	-1.232372	0.784956	-0.937093
459	-1.240701	2.071676	-1.246515	-1.034312	-1.174640	-1.034312

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	fractal
40	-0.317568	-0.305278	-0.051820	0.150716	
212	-0.652519	0.229655	0.682979	-2.024902	
287	-0.887048	-0.736197	-0.927188	-0.956489	
314	-1.122404	-1.304683	-1.743529	0.389937	
459	-0.911200	-0.960044	-1.003254	-0.937093	

5 Creat un modelo

```
In [9]: RANDOM_STATE = 42*42*42
```

```
# Generate a -1, 1 version of Y
y = Y.replace(['B', 'M'], [1, -1])

# Calculate the ratio of B to M samples (aka contamination)
c = Y.value_counts()
ratio = c[1]/(c[0]+c[1])
```

```
In [10]: from sklearn.model_selection import GridSearchCV
         from sklearn.metrics import classification_report
```

```
In [11]: tuned_parameters = [{'n_estimators': [10, 100, 200, 300, 400, 500],
                             'max_samples': ['auto', 0.5, 0.8, 1.0],
                             'max_features': [0.1, 0.5, 0.8, 1.0]}]

clf = GridSearchCV(IsolationForest(behaviour="new", random_state=RANDOM_STATE, contamination=0.37258347978910367),
                  tuned_parameters, cv=5, scoring='f1', n_jobs=-1)

clf.fit(nX, y)
```

```
/Users/manuel.pasieka/anaconda3/envs/py3/lib/python3.6/site-packages/sklearn/model_selection/_search.py:144: DeprecationWarning:
  The parameter 'behaviour' is deprecated in version 0.22 and will be removed in a future version. Use 'bootstrap' instead.
```

```
Out[11]: GridSearchCV(cv=5, error_score='raise-deprecating',
                    estimator=IsolationForest(behaviour='new', bootstrap=False,
                    contamination=0.37258347978910367, max_features=1.0,
                    max_samples='auto', n_estimators=100, n_jobs=None,
                    random_state=74088, verbose=0),
                    fit_params=None, iid='warn', n_jobs=-1,
                    param_grid=[{'n_estimators': [10, 100, 200, 300, 400, 500], 'max_samples': ['auto', 0.5, 0.8, 1.0]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                    scoring='f1', verbose=0)
```

```
In [12]: clf.best_estimator_
```

```
Out[12]: IsolationForest(behaviour='new', bootstrap=False,
                    contamination=0.37258347978910367, max_features=0.1,
```

```
max_samples=1.0, n_estimators=100, n_jobs=None, random_state=74088,
verbose=0)
```

```
In [13]: best = clf.best_estimator_
p = best.predict(nX)
print(classification_report(y, p))
```

	precision	recall	f1-score	support
-1	0.68	0.68	0.68	212
1	0.81	0.81	0.81	357
micro avg	0.76	0.76	0.76	569
macro avg	0.74	0.74	0.74	569
weighted avg	0.76	0.76	0.76	569

```
In [14]: # Confusion matrix
from sklearn.metrics import confusion_matrix
sns.heatmap(confusion_matrix(y, p), annot=True, fmt='2d',
            xticklabels=['B', 'M'], yticklabels=['B', 'M'])
plt.show()
```

